



Auto-Encoding Variational Bayes

Kingma, D. P. & Welling, M. (2014)

Variational Graph Auto-Encoders

Kipf, T., & Welling, M. (2016)

24.01.31 (수)

발표자 : 윤민석

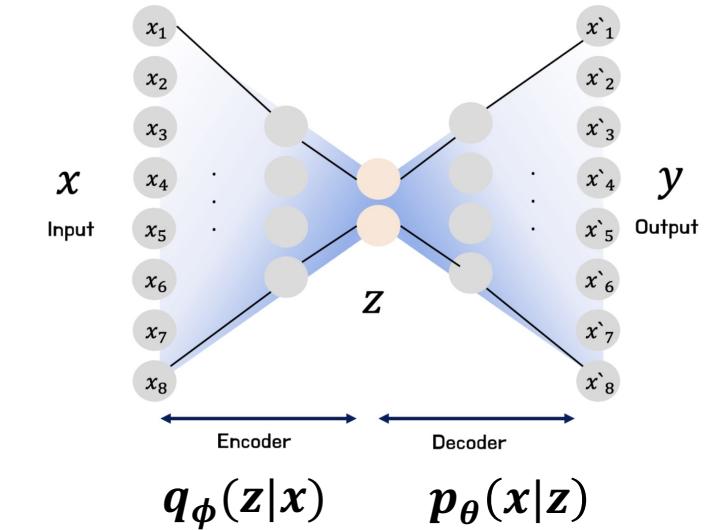
Contents

- ✓ Background
- ✓ Paper #1 (Auto-Encoding Variational Bayes)
 - ✓ Methodology
 - ✓ Code
- ✓ Paper #2 (Variational Graph Auto-Encoders)
 - ✓ Methodology
 - ✓ Code
- ✓ Takeaway



- Autoencoder

- Unsupervised Learning
- Encoder(Recognition network) : Input data를 Latent vector로 변환
- Decoder(Generative network) : Latent vector를 output data로 변환



Background

Background

Paper #1

Paper #2

Takeaway

- MLE & MAP

- Maximum Likelihood Estimation (MLE)

$$\theta_{MLE} = \arg \max_{\theta} \log P(X|\theta)$$

$$= \arg \max_{\theta} \log \prod_i P(x_i|\theta)$$

$$= \arg \max_{\theta} \sum_i \log P(x_i|\theta)$$

- Maximum A posterior Probability (MAP)

$$\theta_{MAP} = \arg \max_{\theta} P(X|\theta)P(\theta)$$

$$P(\theta|X) = \frac{P(X|\theta)p(\theta)}{P(X)} \propto P(X|\theta)p(\theta)$$

$$= \arg \max_{\theta} \log P(X|\theta) + \log P(\theta)$$

$$= \arg \max_{\theta} \log \prod_i P(x_i|\theta) + \log P(\theta)$$

$$= \arg \max_{\theta} \sum_i \log P(x_i|\theta) + \log P(\theta)$$

Background

Background

Paper #1

Paper #2

Takeaway

- MLE
 - Negative Log-likelihood

$$\theta_{MLE} = \arg \max_{\theta} \log P(X|\theta)$$

$$= \arg \max_{\theta} \log \prod_i P(x_i|\theta)$$

$$= \arg \max_{\theta} \sum_i \log P(x_i|\theta)$$



$$\hat{\theta} = \operatorname{argmin}_{\theta \in \Theta} \sum_{i=1}^N -\log P(y_i|x_i; \theta)$$

If Gaussian

If Bernoulli

$$\ln p(\mathbf{X}|\boldsymbol{\mu}, \Sigma) = -\frac{ND}{2} \ln(2\pi) - \frac{N}{2} \ln |\Sigma| - \frac{1}{2} \sum_{n=1}^N (\mathbf{x}_n - \boldsymbol{\mu})^\top \Sigma^{-1} (\mathbf{x}_n - \boldsymbol{\mu})$$

$$-\sum_{i=1}^N \log P(y_i|x_i; \theta) = -\sum_{i=1}^N \mathbf{y}_i^T \cdot \log \hat{\mathbf{y}}_i$$

- KL Divergence

- 두 확률분포의 차이를 계산하는 함수

$$D_{\text{KL}}(P \parallel Q) = \sum_{x \in \mathcal{X}} P(x) \log \left(\frac{P(x)}{Q(x)} \right)$$

$$\begin{aligned} KL(p, q) &= - \int p(x) \log q(x) dx + \int p(x) \log p(x) dx \\ &= \frac{1}{2} \log(2\pi\sigma_2^2) + \frac{\sigma_1^2 + (\mu_1 - \mu_2)^2}{2\sigma_2^2} - \frac{1}{2}(1 + \log 2\pi\sigma_1^2) \\ &= \log \frac{\sigma_2}{\sigma_1} + \frac{\sigma_1^2 + (\mu_1 - \mu_2)^2}{2\sigma_2^2} - \frac{1}{2} \end{aligned}$$

- Monte Carlo simulation

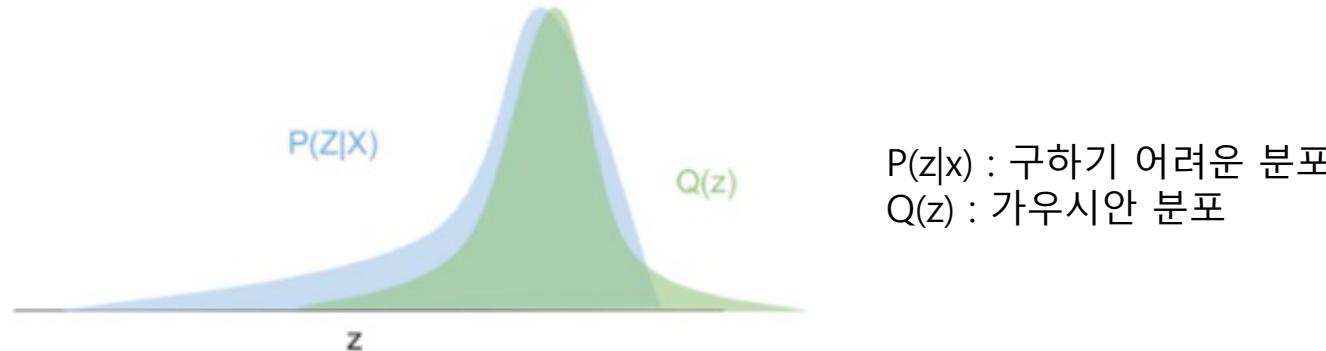
- 대수의 법칙에 의해 모평균으로 수렴

$$E[H(x)] \approx \frac{1}{N} \sum_{i=1}^N H(x_i)$$

$\frac{1}{N} \sum_{i=1}^N H(x_i)$: Monte Carlo Simulation ↗

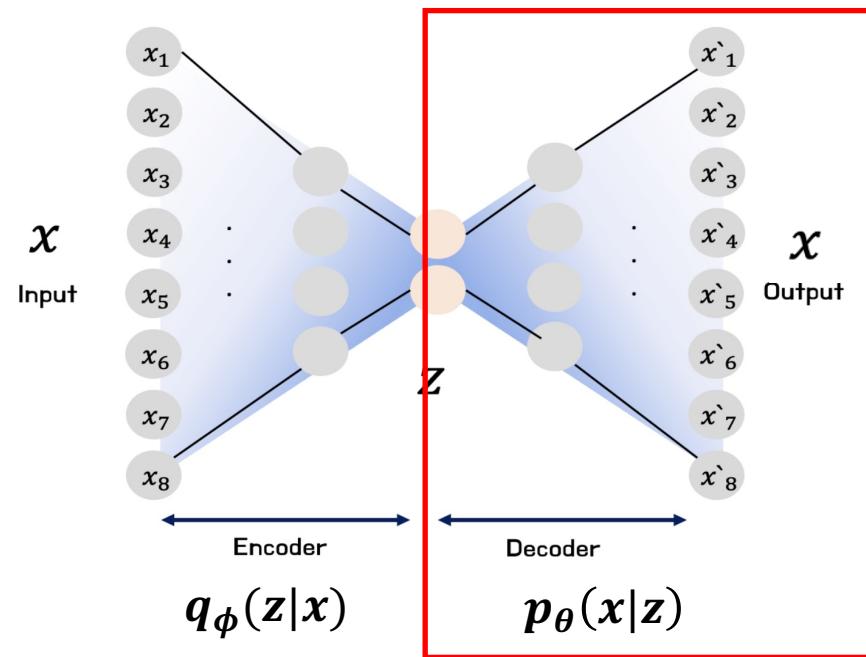
▪ Variational Inference

- 추론 목적: Likelihood 계산 + Posterior($P(z|x)$) 구하기
- 대부분의 경우 $P(z|x)$ 의 정확한 분포를 모름
- 사후확률(Posterior) 분포 $p(z|x)$ 를 다루기 쉬운 확률분포 $q(z)$ 로 근사



Motivation

- How can we perform efficient approximate inference and learning with directed probabilistic models whose continuous latent variables and/or parameters **have intractable posterior distributions?**



모든 경우를 적분하는 것이 어려움

$$p_\theta(x) = \int p_\theta(x|z)p_\theta(z) dz$$

분포 가정 분포 가정

$p_\theta(z|x)$ 대신 $q_\phi(z|x)$ 를 쓰자

Variational Inference

$$p_\theta(z|x) = \frac{p_\theta(x|z)p_\theta(z)}{p_\theta(x)}$$

Intractable Posterior 구하기 어려움

- Directed Graphical model

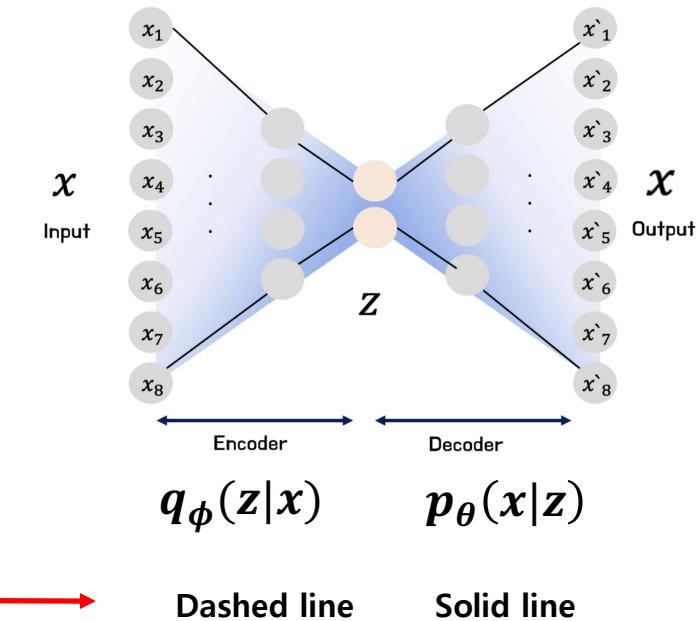
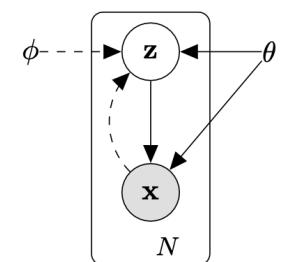
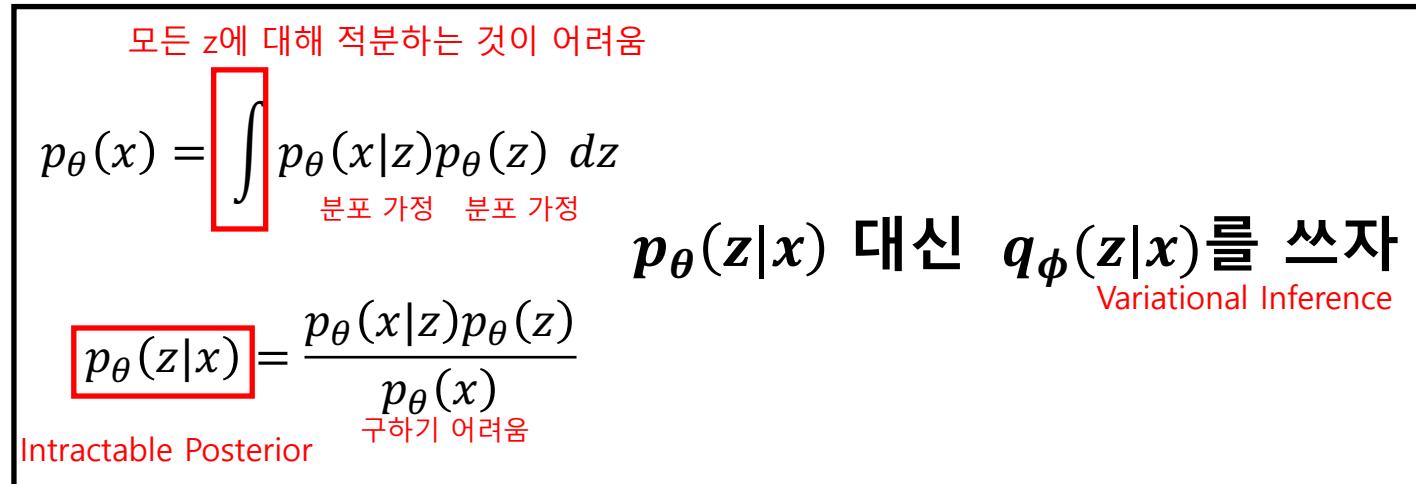


Figure 1: The type of directed graphical model under consideration. Solid lines denote the generative model $p_{\theta}(z)p_{\theta}(x|z)$, dashed lines denote the variational approximation $q_{\phi}(z|x)$ to the intractable posterior $p_{\theta}(z|x)$. The variational parameters ϕ are learned jointly with the generative model parameters θ .

- Variational Inference (Evidence Lower BOund)

$$\log p_{\theta}(\mathbf{x}^{(i)}) = D_{KL}(q_{\phi}(\mathbf{z}|\mathbf{x}^{(i)})||p_{\theta}(\mathbf{z}|\mathbf{x}^{(i)})) + \mathcal{L}(\theta, \phi; \mathbf{x}^{(i)}) \quad (1)$$

$$\log p_{\theta}(\mathbf{x}^{(i)}) \geq \mathcal{L}(\theta, \phi; \mathbf{x}^{(i)}) = \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} [-\log q_{\phi}(\mathbf{z}|\mathbf{x}) + \log p_{\theta}(\mathbf{x}, \mathbf{z})] \quad (2)$$

$$\log p_{\theta}(x) = \log p_{\theta}(x) \int q_{\phi}(z|x) dz = \int \log p_{\theta}(x) q_{\phi}(z|x) dz$$

$$= \int \log \left[\frac{p_{\theta}(x,z)}{p_{\theta}(z|x)} \right] q_{\phi}(z|x) dz$$

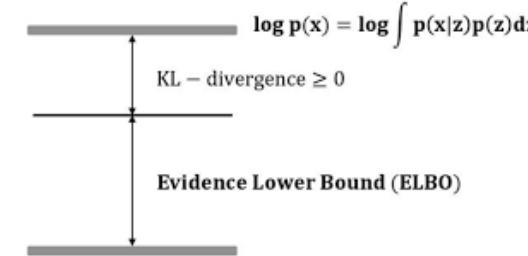
$$= \int \log \left[\frac{p_{\theta}(x,z)}{q_{\phi}(z|x)} \frac{q_{\phi}(z|x)}{p_{\theta}(z|x)} \right] q_{\phi}(z|x) dz$$

$$= \int [\log \left(\frac{p_{\theta}(x,z)}{q_{\phi}(z|x)} \right) + \log \left(\frac{q_{\phi}(z|x)}{p_{\theta}(z|x)} \right)] q_{\phi}(z|x) dz$$

$$= \int \log \frac{p_{\theta}(x,z)p_{\theta}}{q_{\phi}(z|x)} q_{\phi}(z|x) dz + \int \log \frac{q_{\phi}(z|x)}{p_{\theta}(z|x)} q_{\phi}(z|x) dz$$

$$= \mathcal{L}(\theta, \phi; x) + D_{KL}(q_{\phi}(z|x)||p_{\theta}(z|x)) \geq \mathcal{L}(\theta, \phi; x)$$

Intractable Lower bound



- Variational Inference (Evidence Lower BOund)

$$\mathcal{L}(\boldsymbol{\theta}, \boldsymbol{\phi}; \mathbf{x}^{(i)}) = -D_{KL}(q_{\boldsymbol{\phi}}(\mathbf{z}|\mathbf{x}^{(i)})||p_{\boldsymbol{\theta}}(\mathbf{z})) + \mathbb{E}_{q_{\boldsymbol{\phi}}(\mathbf{z}|\mathbf{x}^{(i)})} \left[\log p_{\boldsymbol{\theta}}(\mathbf{x}^{(i)}|\mathbf{z}) \right] \quad (3)$$

$$\mathcal{L}(\boldsymbol{\theta}, \boldsymbol{\phi}; \mathbf{x}) = \int \log \frac{p_{\boldsymbol{\theta}}(\mathbf{x}|z)}{q_{\boldsymbol{\phi}}(z|x)} q_{\boldsymbol{\phi}}(z|x) dz = \int [\log \frac{p_{\boldsymbol{\theta}}(x|z)p(z)}{q_{\boldsymbol{\phi}}(z|x)}] q_{\boldsymbol{\phi}}(z|x) dz$$

Lower bound

$$= \int \log p_{\boldsymbol{\theta}}(x|z) q_{\boldsymbol{\phi}}(z|x) dz + \int \log \frac{p_{\boldsymbol{\theta}}(x)}{q_{\boldsymbol{\phi}}(z|x)} q_{\boldsymbol{\phi}}(z|x) dz$$

$$= \mathbb{E}_{z \sim q_{\boldsymbol{\phi}}(z|x)} [\log p_{\boldsymbol{\theta}}(x|z)] - \int \log \frac{q_{\boldsymbol{\phi}}(z|x)}{p_{\boldsymbol{\theta}}(x)} q_{\boldsymbol{\phi}}(z|x) dz$$

$$= \mathbb{E}_{z \sim q_{\boldsymbol{\phi}}(z|x)} [\log p_{\boldsymbol{\theta}}(x|z)] - D_{KL}(q_{\boldsymbol{\phi}}(z|x)||p_{\boldsymbol{\theta}}(z))$$

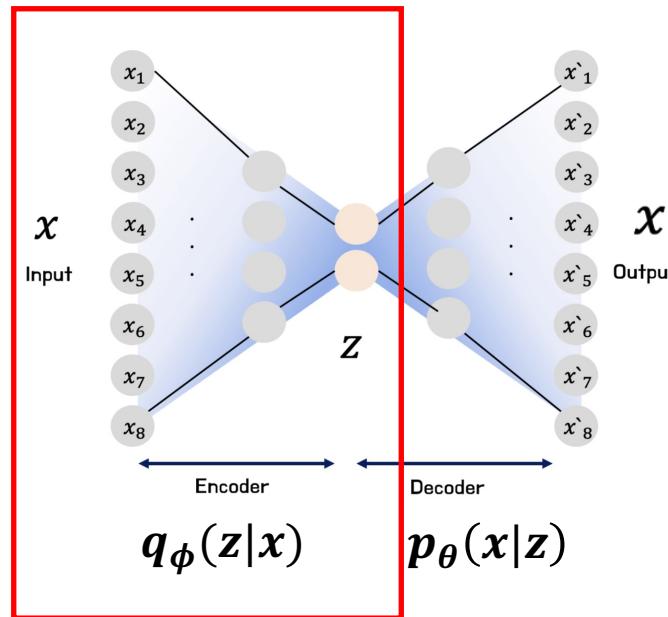
Reconstruction

Regularization

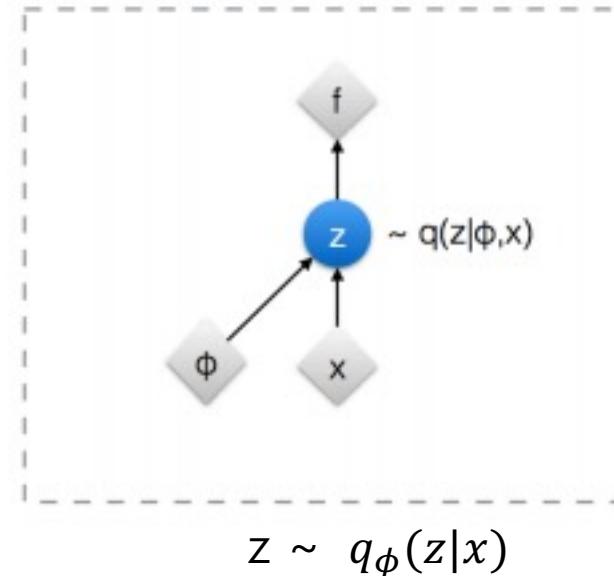
Paper #1 (VAE)

[Background](#)[Paper #1](#)[Paper #2](#)[Takeaway](#)

- Reparametrization Trick

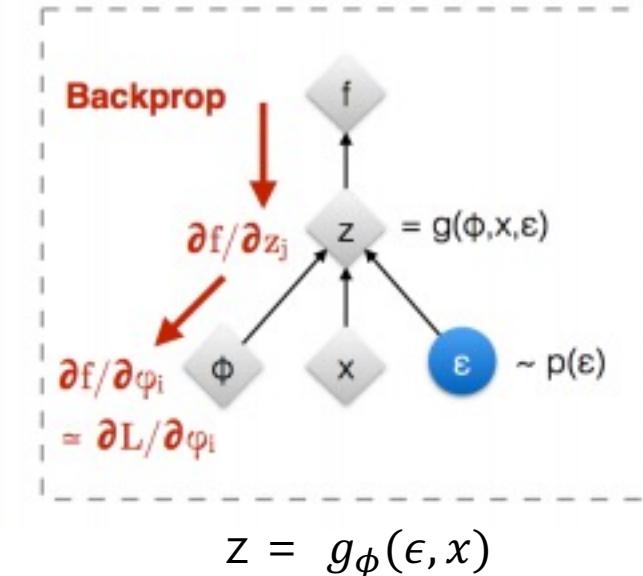


Original form



High variance

Reparameterised form



$z = \mu + \sigma\epsilon$

■ ELBO Structure

- $\mathcal{L}(\theta, \phi; x) = -D_{KL}(q_\phi(z|x) || p_\theta(z)) + \mathbb{E}_{z \sim q_\phi(z|x)} [\log p_\theta(x|z)]$

Monte Carlo simulation

$$\tilde{\mathcal{L}}^B(\boldsymbol{\theta}, \boldsymbol{\phi}; \mathbf{x}^{(i)}) = -D_{KL}(q_\phi(\mathbf{z}|\mathbf{x}^{(i)}) || p_\theta(\mathbf{z})) + \frac{1}{L} \sum_{l=1}^L (\log p_\theta(\mathbf{x}^{(i)}|\mathbf{z}^{(i,l)}))$$

where $\mathbf{z}^{(i,l)} = g_\phi(\boldsymbol{\epsilon}^{(i,l)}, \mathbf{x}^{(i)})$ and $\boldsymbol{\epsilon}^{(l)} \sim p(\boldsymbol{\epsilon})$ (7)

Reparametrization Trick

$$\mathcal{L}(\boldsymbol{\theta}, \boldsymbol{\phi}; \mathbf{x}^{(i)}) \simeq \frac{1}{2} \sum_{j=1}^J \left(1 + \log((\sigma_j^{(i)})^2) - (\mu_j^{(i)})^2 - (\sigma_j^{(i)})^2 \right) + \frac{1}{L} \sum_{l=1}^L \log p_\theta(\mathbf{x}^{(i)}|\mathbf{z}^{(i,l)})$$

where $\mathbf{z}^{(i,l)} = \boldsymbol{\mu}^{(i)} + \boldsymbol{\sigma}^{(i)} \odot \boldsymbol{\epsilon}^{(l)}$ and $\boldsymbol{\epsilon}^{(l)} \sim \mathcal{N}(0, \mathbf{I})$ (10)

Paper #1 (VAE)

Background

Paper #1

Paper #2

Takeaway

■ Code

```
● ○ ●  
1 import torch  
2 import torch.nn as nn  
3 import torch.nn.functional as F  
4 import torch.optim  
5 from torchvision import datasets, transforms  
6 from torchvision.utils import save_image  
7 import matplotlib.pyplot as plt  
8  
9 device = 'cpu'  
10 # 데이터셋 불러오기  
11  
12 train_dataset = datasets.MNIST(root='./mnist_data/', train=True, transform=transforms.ToTensor(), download=True)  
13 test_dataset = datasets.MNIST(root='./mnist_data/', train=False, transform=transforms.ToTensor(), download=False)  
14 train_loader = torch.utils.data.DataLoader(dataset=train_dataset, batch_size=100, shuffle=True)  
15 test_loader = torch.utils.data.DataLoader(dataset=test_dataset, batch_size=100, shuffle=False)
```

```
● ○ ●  
1 def loss_function(recon_x, x, average, log_variance):  
2     # Minimize 상태로 만들기  
3     binary_cross_entropy = F.binary_cross_entropy(recon_x, x.view(-1, 784), reduction='sum')  
4     KL_divergence = - 0.5 * torch.sum(1 + log_variance - average ** 2 - log_variance.exp())  
5     return binary_cross_entropy + KL_divergence
```

```
● ○ ●  
1 class VAE(nn.Module):  
2     def __init__(self, x_dim, h_dim1, h_dim2, z_dim):  
3         super(VAE, self).__init__()  
4  
5         self.fc1 = nn.Linear(x_dim, h_dim1)  
6         self.fc2 = nn.Linear(h_dim1, h_dim2)  
7         self.fc_average = nn.Linear(h_dim2, z_dim)  
8         self.fc_log_variance = nn.Linear(h_dim2, z_dim)  
9  
10        self.fc3 = nn.Linear(z_dim, h_dim2)  
11        self.fc4 = nn.Linear(h_dim2, h_dim1)  
12        self.fc5 = nn.Linear(h_dim1, x_dim)  
13  
14    def encoder(self, x):  
15        x = F.relu(self.fc1(x))  
16        x = F.relu(self.fc2(x))  
17        return self.fc_average(x), self.fc_log_variance(x)  
18  
19    def reparameterization(self, average, log_variance):  
20        std = torch.exp(0.5 * log_variance)  
21        eps = torch.randn_like(std)  
22        return average + eps * std  
23  
24    def decoder(self, z):  
25        z = F.relu(self.fc3(z))  
26        z = F.relu(self.fc4(z))  
27        return torch.sigmoid(self.fc5(z))  
28  
29    def forward(self, x):  
30        average, log_variance = self.encoder(x.view(-1, 784))  
31        z = self.reparameterization(average, log_variance)  
32        return self.decoder(z), average, log_variance
```

Paper #1 (VAE)

Background

Paper #1

Paper #2

Takeaway

■ Code

```
● ● ●  
1 train_loss_list = []  
2 test_loss_list = []  
3  
4 def train(epoch):  
5     model.train()  
6     train_loss = 0  
7     for data, _ in train_loader:  
8         data = data.to(device)  
9         optimizer.zero_grad()  
10        x_reconstruct, average, log_variance = model(data)  
11        loss = loss_function(x_reconstruct, data, average, log_variance)  
12        loss.backward()  
13        train_loss += loss.item()  
14        optimizer.step()  
15        print("{}'s avrage train loss : {}".format(epoch, train_loss / len(train_loader.dataset)))  
16        train_loss_list.append(train_loss / len(train_loader.dataset))  
17  
18 def test(epoch):  
19     model.eval()  
20     test_loss = 0  
21     with torch.no_grad():  
22         for data, _ in test_loader:  
23             data = data.to(device)  
24             x_reconstruct, average, log_variance = model(data)  
25             test_loss += loss_function(x_reconstruct, data, average, log_variance).item()  
26             print("{}'s avrage test loss : {}".format(epoch, test_loss / len(test_loader.dataset)))  
27             test_loss_list.append(test_loss / len(test_loader.dataset))
```

```
● ● ●  
1 z_list = [3, 5, 10, 20 ,200]  
2 for z in z_list:  
3     train_loader = torch.utils.data.DataLoader(dataset=train_dataset, batch_size=100, shuffle=True)  
4     test_loader = torch.utils.data.DataLoader(dataset=test_dataset, batch_size=100, shuffle=False)  
5     model = VAE(x_dim = 28 ** 2, h_dim1 = 500, h_dim2 = 250, z_dim = z).to(device)  
6     optimizer = torch.optim.Adagrad(model.parameters(), lr = 0.01)  
7     train_loss_list = []  
8     test_loss_list = []  
9     for epoch in range(1, 21):  
10        train(epoch)  
11        test(epoch)  
12  
13    epochs = range(1, 21)  
14    plt.plot(epochs, train_loss_list, label='Train Loss')  
15    plt.plot(epochs, test_loss_list, label='Test Loss')  
16    plt.ylim(100, 150)  
17    plt.xticks(range(1, 21))  
18    plt.xlabel('Epochs')  
19    plt.ylabel('Loss')  
20    plt.legend()  
21    plt.title('Train and Test Loss Over Epochs')  
22    plt.show()
```

Paper #1 (VAE)

Background

Paper #1

Paper #2

Takeaway

■ Code

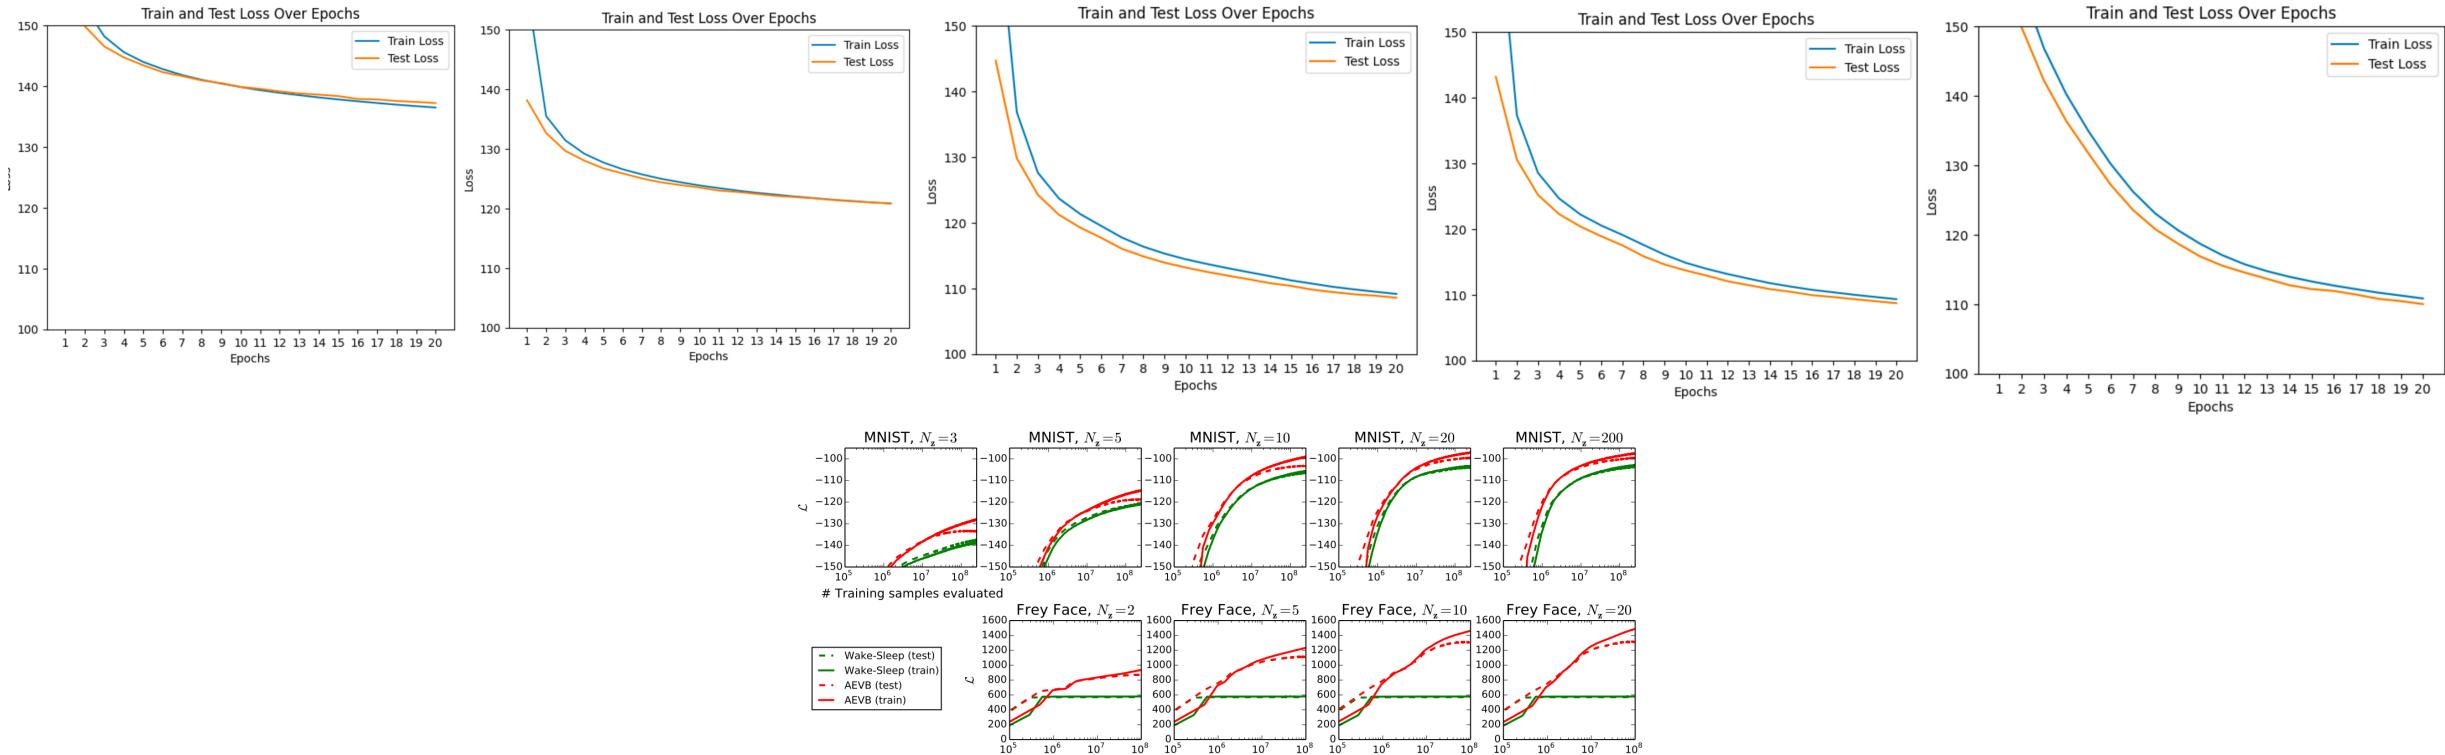


Figure 2: Comparison of our AEVB method to the wake-sleep algorithm, in terms of optimizing the lower bound, for different dimensionality of latent space (N_z). Our method converged considerably faster and reached a better solution in all experiments. Interestingly enough, more latent variables does not result in more overfitting, which is explained by the regularizing effect of the lower bound. Vertical axis: the estimated average variational lower bound per datapoint. The estimator variance was small (< 1) and omitted. Horizontal axis: amount of training points evaluated. Computation took around 20-40 minutes per million training samples with a Intel Xeon CPU running at an effective 40 GFLOPS.

Paper #1 (VAE)

Background

Paper #1

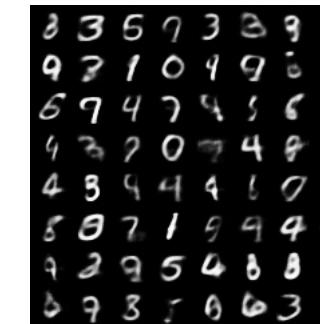
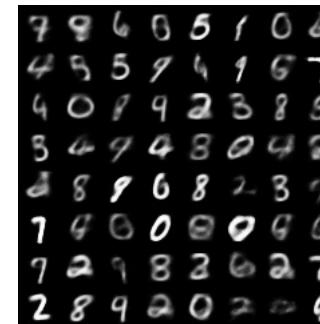
Paper #2

Takeaway

■ Code

```
● ● ●  
1 z_list = [2, 5, 10, 20]  
2 for z in z_list:  
3     train_loader = torch.utils.data.DataLoader(dataset=train_dataset, batch_size=100, shuffle=True)  
4     test_loader = torch.utils.data.DataLoader(dataset=test_dataset, batch_size=100, shuffle=False)  
5     model = VAE(x_dim = 28 ** 2, h_dim1 = 500, h_dim2 = 250, z_dim = z).to(device)  
6     optimizer = torch.optim.Adagrad(model.parameters(), lr = 0.01)  
7     train_loss_list = []  
8     test_loss_list = []  
9     for epoch in range(1, 21):  
10         train(epoch)  
11         test(epoch)  
12  
13     with torch.no_grad():  
14         z_random = torch.randn(64, z).to(device)  
15         sample = model.decoder(z_random).to(device)  
16  
17     save_image(sample.view(64, 1, 28, 28), './samples' + '{}.png'.format(z))
```

- Code



(a) 2-D latent space

(b) 5-D latent space

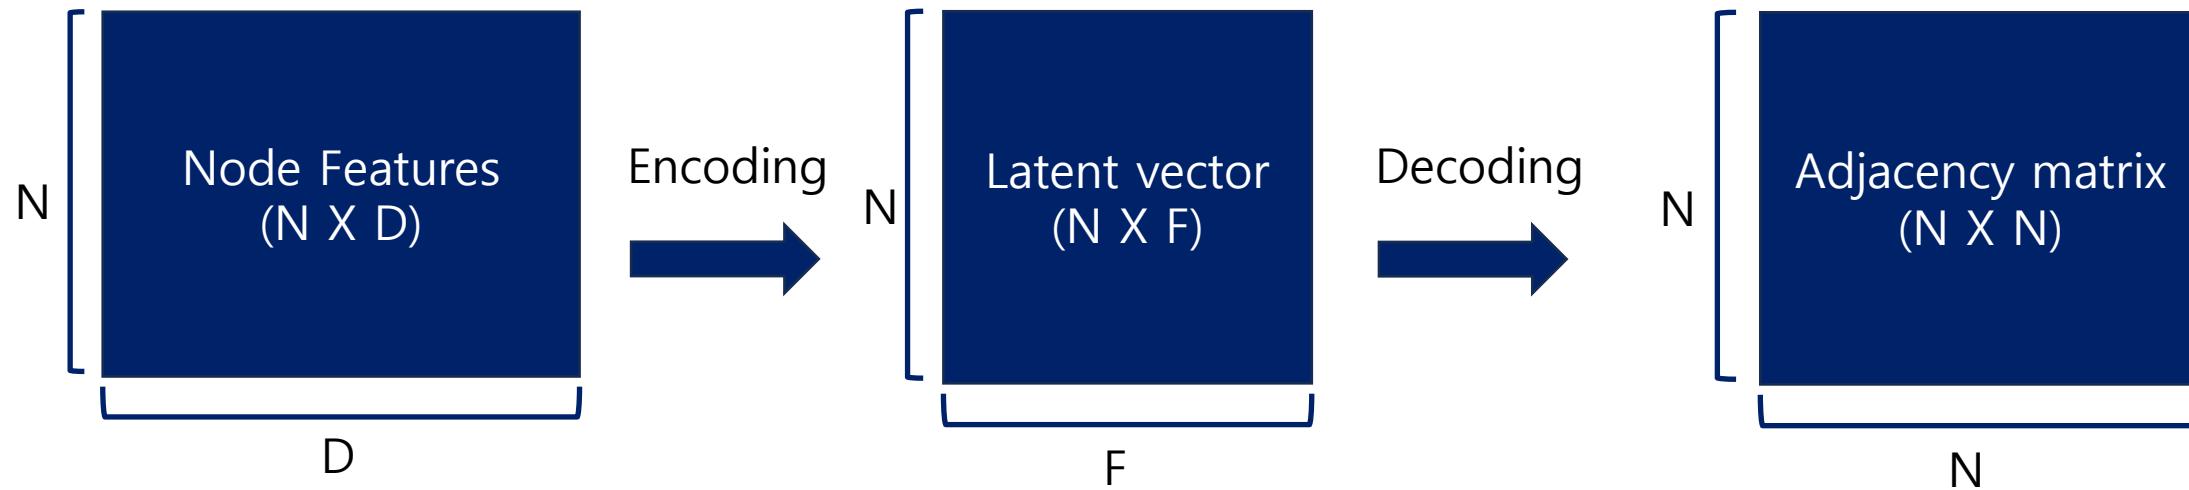
(c) 10-D latent space

(d) 20-D latent space

Figure 5: Random samples from learned generative models of MNIST for different dimensionalities of latent space.

Motivation

- We introduce the variational graph autoencoder (VGAE), a framework for unsupervised learning on graph-structured data based on the variational auto-encoder (VAE).
- Use of latent variables
- Link prediction problem



Paper #2 (VGAE)

[Background](#)[Paper #1](#)[Paper #2](#)[Takeaway](#)

Encoder

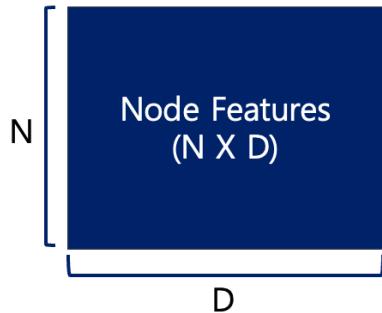
- Using GCN
- Reparametrization trick

$$q(\mathbf{Z} | \mathbf{X}, \mathbf{A}) = \prod_{i=1}^N q(\mathbf{z}_i | \mathbf{X}, \mathbf{A}), \text{ with } q(\mathbf{z}_i | \mathbf{X}, \mathbf{A}) = \mathcal{N}(\mathbf{z}_i | \boldsymbol{\mu}_i, \text{diag}(\boldsymbol{\sigma}_i^2)).$$

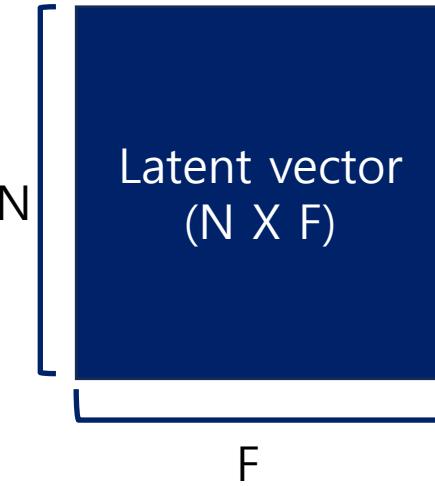
$$\text{GCN}(\mathbf{X}, \mathbf{A}) = \tilde{\mathbf{A}} \text{ReLU}(\tilde{\mathbf{A}} \mathbf{X} \mathbf{W}_0) \mathbf{W}_1$$

$$\boldsymbol{\mu} = \text{GCN}_{\boldsymbol{\mu}}(\mathbf{X}, \mathbf{A})$$

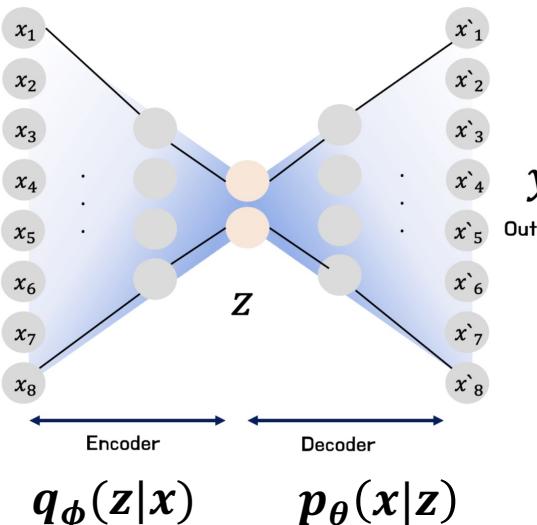
$$\log \boldsymbol{\sigma} = \text{GCN}_{\boldsymbol{\sigma}}(\mathbf{X}, \mathbf{A})$$



Encoding



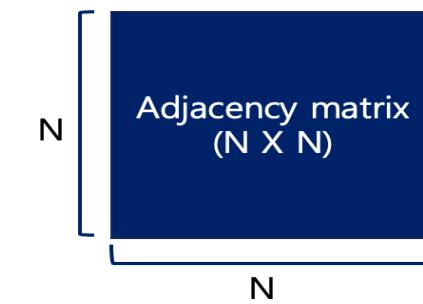
\mathbf{x}
Input



Decoder

- Using inner product

$$p(\mathbf{A} | \mathbf{Z}) = \prod_{i=1}^N \prod_{j=1}^N p(A_{ij} | \mathbf{z}_i, \mathbf{z}_j), \text{ with } p(A_{ij} = 1 | \mathbf{z}_i, \mathbf{z}_j) = \sigma(\mathbf{z}_i^\top \mathbf{z}_j).$$



Decoding



- Learning

- Optimize function

$$\mathcal{L} = \mathbb{E}_{q(\mathbf{Z}|\mathbf{X}, \mathbf{A})} [\log p(\mathbf{A}|\mathbf{Z})] - \text{KL}[q(\mathbf{Z}|\mathbf{X}, \mathbf{A}) || p(\mathbf{Z})]$$

- Featureless approach (VGAE*)
 - Replace X with the identity matrix in the GCN

$$\text{GCN}(\mathbf{X}, \mathbf{A}) = \tilde{\mathbf{A}} \text{ReLU}(\tilde{\mathbf{A}}\mathbf{X}\mathbf{W}_0)\mathbf{W}_1$$

$$\boldsymbol{\mu} = \text{GCN}_{\boldsymbol{\mu}}(\mathbf{X}, \mathbf{A})$$

$$\log \boldsymbol{\sigma} = \text{GCN}_{\boldsymbol{\sigma}}(\mathbf{X}, \mathbf{A})$$

- GAE

- A non-probabilistic variant of the VGAE model
- Only use reconstruction error

$$\hat{\mathbf{A}} = \sigma(\mathbf{Z}\mathbf{Z}^\top), \text{ with } \mathbf{Z} = \text{GCN}(\mathbf{X}, \mathbf{A})$$

$$\mathcal{L} = \mathbb{E}_{q(\mathbf{z}|\mathbf{x}, \mathbf{a})} [\log p(\mathbf{a} | \mathbf{z})] - \text{KL}[q(\mathbf{z} | \mathbf{x}, \mathbf{a}) || p(\mathbf{z})]$$


- Featureless approach (GAE*)
 - Replace X with the identity matrix in the GCN

Experiments

- To learn meaningful latent embeddings on a link prediction task on several popular citation network datasets.
- Split Train / Validation / Test set (**disjoint**)
- Use negative sampling

Table 1: Link prediction task in citation networks. See [1] for dataset details.

Method	Cora		Citeseer		Pubmed	
	AUC	AP	AUC	AP	AUC	AP
SC [5]	84.6 \pm 0.01	88.5 \pm 0.00	80.5 \pm 0.01	85.0 \pm 0.01	84.2 \pm 0.02	87.8 \pm 0.01
DW [6]	83.1 \pm 0.01	85.0 \pm 0.00	80.5 \pm 0.02	83.6 \pm 0.01	84.4 \pm 0.00	84.1 \pm 0.00
GAE*	84.3 \pm 0.02	88.1 \pm 0.01	78.7 \pm 0.02	84.1 \pm 0.02	82.2 \pm 0.01	87.4 \pm 0.00
VGAE*	84.0 \pm 0.02	87.7 \pm 0.01	78.9 \pm 0.03	84.1 \pm 0.02	82.7 \pm 0.01	87.5 \pm 0.01
GAE	91.0 \pm 0.02	92.0 \pm 0.03	89.5 \pm 0.04	89.9 \pm 0.05	96.4 \pm 0.00	96.5 \pm 0.00
VGAE	91.4 \pm 0.01	92.6 \pm 0.01	90.8 \pm 0.02	92.0 \pm 0.02	94.4 \pm 0.02	94.7 \pm 0.02

Paper #2 (VGAE)

Background

Paper #1

Paper #2

Takeaway

■ Code (Failure Case)



```
1 import numpy as np
2 import torch
3 import torch.nn as nn
4 import torch.nn.functional as F
5 import torch.optim
6 import dgl
7 from dgl.nn import GraphConv
8 from torch_geometric.datasets import Planetoid
9 from torch_geometric.transforms import RandomLinkSplit
10
11 device = 'cpu'
12
13 # 데이터 불러오기
14 dataset = Planetoid(root='/tmp/Cora', name='Cora')
15 data = dataset[0]
16
17 # Train / Validation / Test split
18 transform = RandomLinkSplit(is_undirected=True, num_val = 0.05, num_test = 0.1)
19 train_data, val_data, test_data = transform(data)
20
21 num_nodes = data.num_nodes
```



```
1 # 인접행렬 만들기
2 train_edge_index = train_data.edge_index
3 train_adjacency_matrix = torch.zeros((num_nodes, num_nodes))
4 train_adjacency_matrix[train_edge_index[0], train_edge_index[1]] = 1
5 train_adjacency_matrix = (train_adjacency_matrix + np.identity(2708)).to(torch.float32)
6
7 test_edge_index = test_data.edge_index
8 test_adjacency_matrix = torch.zeros((num_nodes, num_nodes))
9 test_adjacency_matrix[test_edge_index[0], test_edge_index[1]] = 1
10 test_adjacency_matrix = (test_adjacency_matrix + np.identity(2708)).to(torch.float32)
11
12 node_features = data.x
13
14
15 # 그래프 데이터 만들기
16 train_edge_index = train_edge_index.tolist()
17 train_a = torch.tensor(train_edge_index[0])
18 train_b = torch.tensor(train_edge_index[1])
19 train_graph = dgl.graph((train_a, train_b))
20 train_graph = dgl.add_self_loop(train_graph)
21
22 test_edge_index = test_edge_index.tolist()
23 test_a = torch.tensor(test_edge_index[0])
24 test_b = torch.tensor(test_edge_index[1])
25 test_graph = dgl.graph((test_a, test_b))
26 test_graph = dgl.add_self_loop(test_graph)
27
28 train_graph = train_graph.to(device)
29 test_graph = test_graph.to(device)
```

Paper #2 (VGAE)

Background

Paper #1

Paper #2

Takeaway

■ Code (Failure Case)

```
● ● ●  
1 class VGAE(nn.Module):  
2     def __init__(self, in_dim = 1433, hidden1_dim = 32, z = 16):  
3         super().__init__()  
4         self.graph_conv1 = GraphConv(in_dim, hidden1_dim, activation = F.relu, allow_zero_in_degree = True)  
5         self.graph_conv_average = GraphConv(hidden1_dim, z, activation = lambda x: x, allow_zero_in_degree = True)  
6         self.graph_conv_log_variance = GraphConv(hidden1_dim, z, activation = lambda x: x, allow_zero_in_degree = True)  
7  
8     def encoder(self, adj, features):  
9         z = self.graph_conv1(adj, features)  
10        average = self.graph_conv_average(adj, z)  
11        log_variance = self.graph_conv_log_variance(adj, z)  
12        return average, log_variance  
13  
14    def reparameterization(self, average, log_variance):  
15        std = torch.exp(0.5 * log_variance)  
16        eps = torch.randn_like(std)  
17        return average + std * eps  
18  
19    def decoder(self, z):  
20        new_adj = torch.sigmoid(torch.matmul(z, z.t()))  
21        return new_adj  
22  
23    def forward(self, adj, features):  
24        average, log_variance = self.encoder(adj, features)  
25        z = self.reparameterization(average, log_variance)  
26        new_adj = self.decoder(z)  
27        return new_adj, average, log_variance
```

```
● ● ●  
1 def loss_function(new_adj, adj, average, log_variance):  
2     # Minimize 상태로 만들기  
3     binary_cross_entropy = F.binary_cross_entropy(new_adj, adj, reduction='sum')  
4     KL_divergence = - 0.5 * torch.sum(1 + log_variance - average ** 2 - log_variance.exp())  
5     return binary_cross_entropy + KL_divergence
```

```
● ● ●  
1 train_loss_list = []  
2 test_loss_list = []  
3 model = VGAE().to(device)  
4 optimizer = torch.optim.Adam(model.parameters(), lr = 0.03)
```

Paper #2 (VGAE)

Background

Paper #1

Paper #2

Takeaway

■ Code (Failure Case)

```
● ● ●  
1 def train():  
2     with train_graph.local_scope():  
3         model.train()  
4         optimizer.zero_grad()  
5         new_adj, average, log_variance = model(train_graph, node_features)  
6         loss = loss_function(new_adj, train_adjacency_matrix, average, log_variance)  
7         loss.backward()  
8         optimizer.step()  
9         train_loss_list.append(loss.item())  
10        print("Train Loss: {}".format(loss.item()))  
11  
12  
13 def test():  
14     model.eval()  
15     with test_graph.local_scope():  
16         with torch.no_grad():  
17             new_adj, average, log_variance = model(test_graph, node_features)  
18             loss = loss_function(new_adj, test_adjacency_matrix, average, log_variance)  
19             train_loss_list.append(loss.item())  
20             print("test Loss: {}".format(loss.item()))  
21     return new_adj
```

```
● ● ●  
1 for epoch in range(1,201):  
2     train()  
3     test()
```

```
Train Loss: 5259356.0  
test Loss: 5216732.0  
Train Loss: 5213807.0  
test Loss: 5201317.5  
Train Loss: 5201628.0  
test Loss: 5217998.0  
Train Loss: 5219848.0  
test Loss: 5227190.5  
Train Loss: 5227084.5  
  
show more (open the raw output data in a text editor) ...  
  
test Loss: 5157293.5  
Train Loss: 5156840.0  
test Loss: 5156893.5  
Train Loss: 5156632.5  
test Loss: 5156985.5
```

```
● ● ●  
1 new_adj = np.array(test())  
2 binary_adj = (new_adj >= 0.5).astype(int)  
3  
4 test_adjacency_matrix = np.array(test_adjacency_matrix)  
5  
6 from sklearn.metrics import roc_auc_score, precision_score  
7 auc_score = roc_auc_score(binary_adj.flatten(), test_adjacency_matrix.flatten())  
8 precision = precision_score(binary_adj.flatten(), test_adjacency_matrix.flatten(), average='binary')  
9 print("AUC Score:", auc_score)  
10 print("Precision:", precision)
```

AUC Score: 0.5004482530894087
Precision: 0.6376740376740376

Paper #2 (VGAE)

Background

Paper #1

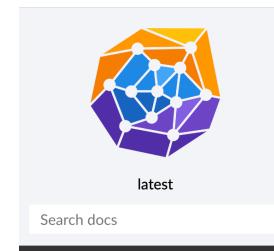
Paper #2

Takeaway

■ Code (Success Case)

```
● ● ●  
1 import numpy as np  
2 import torch  
3 import torch_geometric.transforms as T  
4 from torch_geometric.datasets import Planetoid  
5 from torch_geometric.nn import GCNConv, VGAE  
6  
7 device = 'cpu'  
8 transform = T.Compose([  
9     T.NormalizeFeatures(),  
10    T.RandomLinkSplit(num_val=0.05, num_test=0.1, is_undirected=True, split_labels=True, add_negative_train_samples=False),  
11])  
12  
13 dataset = Planetoid('.', name='Cora', transform=transform)  
14 train_data, val_data, test_data = dataset[0]  
15 train_data, val_data, test_data = train_data.to(device), val_data.to(device), test_data.to(device)
```

```
● ● ●  
1 class Encoder(torch.nn.Module):  
2     def __init__(self, dim_in, hidden_dim = 32, dim_out = 16):  
3         super().__init__()  
4         self.fc1 = GCNConv(dim_in, hidden_dim)  
5         self.fc_average = GCNConv(hidden_dim, dim_out)  
6         self.fc_log_variance = GCNConv(hidden_dim, dim_out)  
7  
8     def forward(self, x, edge_index):  
9         x = self.fc1(x, edge_index).relu()  
10        return self.fc_average(x, edge_index), self.fc_log_variance(x, edge_index)  
11  
12 model = VGAE(Encoder(dataset.num_features, 32, 16)).to(device)  
13 optimizer = torch.optim.Adam(model.parameters(), lr=0.01)
```



Module code / models.autoencoder

Source code for
torch_geometric.nn.models.autoencoder

Paper #2 (VGAE)

Background

Paper #1

Paper #2

Takeaway

■ Code (Success Case)

```
● ● ●  
1 def train():  
2     model.train()  
3     optimizer.zero_grad()  
4     z = model.encode(train_data.x, train_data.edge_index)  
5     loss = model.recon_loss(z, train_data.pos_edge_label_index) + (1 / train_data.num_nodes) * model.kl_loss()  
6     loss.backward()  
7     optimizer.step()  
8     return float(loss)  
9  
10 @torch.no_grad()  
11 def test(data):  
12     model.eval()  
13     z = model.encode(data.x, data.edge_index)  
14     return model.test(z, data.pos_edge_label_index, data.neg_edge_label_index)  
15  
16  
17 for epoch in range(201):  
18     loss = train()  
19     test_auc, test_ap = test(test_data)  
20     if epoch % 50 == 0:  
21         print('Train loss : {}'.format(loss))  
22         print('Test AUC : {} | Test AP : {}'.format(test_auc, test_ap))
```

```
def recon_loss(self, z: Tensor, pos_edge_index: Tensor,  
              neg_edge_index: Optional[Tensor] = None) -> Tensor:  
    r"""Given latent variables :obj:`z`, computes the binary cross  
    entropy loss for positive edges :obj:`pos_edge_index` and negative  
    sampled edges.  
    """
```

Args:

```
z (torch.Tensor): The latent space :math:`\mathbf{Z}`.  
pos_edge_index (torch.Tensor): The positive edges to train against.  
neg_edge_index (torch.Tensor, optional): The negative edges to  
train against. If not given, uses negative sampling to  
calculate negative edges. (default: :obj:`None`)  
"""\n    pos_loss = -torch.log(\n        self.decoder(z, pos_edge_index, sigmoid=True) + EPS).mean()  
  
    if neg_edge_index is None:  
        neg_edge_index = negative_sampling(pos_edge_index, z.size(0))  
    neg_loss = -torch.log(1 -  
        self.decoder(z, neg_edge_index, sigmoid=True) +  
        EPS).mean()  
  
    return pos_loss + neg_loss
```

```
def test(self, z: Tensor, pos_edge_index: Tensor,  
        neg_edge_index: Tensor) -> Tuple[Tensor, Tensor]:  
    r"""Given latent variables :obj:`z`, positive edges  
    :obj:`pos_edge_index` and negative edges :obj:`neg_edge_index`,  
    computes area under the ROC curve (AUC) and average precision (AP)  
    scores.  
    """
```

Args:

```
z (torch.Tensor): The latent space :math:`\mathbf{Z}`.  
pos_edge_index (torch.Tensor): The positive edges to evaluate  
against.  
neg_edge_index (torch.Tensor): The negative edges to evaluate  
against.  
"""\n    from sklearn.metrics import average_precision_score, roc_auc_score  
  
    pos_y = z.new_ones(pos_edge_index.size(1))  
    neg_y = z.new_zeros(neg_edge_index.size(1))  
    y = torch.cat([pos_y, neg_y], dim=0)  
  
    pos_pred = self.decoder(z, pos_edge_index, sigmoid=True)  
    neg_pred = self.decoder(z, neg_edge_index, sigmoid=True)  
    pred = torch.cat([pos_pred, neg_pred], dim=0)  
  
    y, pred = y.detach().cpu().numpy(), pred.detach().cpu().numpy()  
  
    return roc_auc_score(y, pred), average_precision_score(y, pred)
```

Paper #2 (VGAE)

Background

Paper #1

Paper #2

Takeaway

■ Code (Success Case)

```
● ● ●  
1 def train():  
2     model.train()  
3     optimizer.zero_grad()  
4     z = model.encode(train_data.x, train_data.edge_index)  
5     loss = model.recon_loss(z, train_data.pos_edge_label_index) + (1 / train_data.num_nodes) * model.kl_loss()  
6     loss.backward()  
7     optimizer.step()  
8     return float(loss)  
9  
10 @torch.no_grad()  
11 def test(data):  
12     model.eval()  
13     z = model.encode(data.x, data.edge_index)  
14     return model.test(z, data.pos_edge_label_index, data.neg_edge_label_index)  
15  
16  
17 for epoch in range(201):  
18     loss = train()  
19     test_auc, test_ap = test(test_data)  
20     if epoch % 50 == 0:  
21         print('Train loss : {}'.format(loss))  
22         print('Test AUC : {} | Test AP : {}'.format(test_auc, test_ap))
```

```
Train loss : 3.5164217948913574  
Test AUC : 0.6655624727702184 | Test AP : 0.69503406826326  
Train loss : 1.3143773078918457  
Test AUC : 0.6430261153858617 | Test AP : 0.6786212637061367  
Train loss : 1.146220326423645  
Test AUC : 0.7295637113877196 | Test AP : 0.7447826307149207  
Train loss : 1.0693825483322144  
Test AUC : 0.7687889993482855 | Test AP : 0.7783319383978209  
Train loss : 0.9703818559646606  
Test AUC : 0.8702080085262971 | Test AP : 0.8629827566623429
```

(0.8641732568057945, 0.8772858317978459)

- Good
 - 전반적인 흐름 이해 (Background, Methodology, ...)
 - VAE 재현성 확인

- Weak
 - Link prediction 실험 설계 이해 부족
 - Failure Case 성공적으로 돌려보기

- Takeaway
 - Theoretical background의 중요성 (Directed graphical model, ELBO, ...)
 - Library 패키지를 사용한 공부
 - GNN 흥미 & Use case..?

Thank you for watching!

- 코드 깃허브 업로드 완료
 - VAE / VGAE_failure / VGAE_success